

# FINAL CENTRAL LP

March 17, 2024

```
[2]: #importing libraries
import numpy as np
import pandas as pd
```

```
[3]: #importing data
df=pd.read_csv("C:\\Users\\hp\\Downloads\\Health_Care_dataset.csv")
df
```

```
[3]:
```

	Age	Heart_disease
0	3.745401	186.756403
1	9.507143	653.040955
2	7.319939	438.165461
3	5.986585	319.056918
4	1.560186	94.475345
..	...	...
95	4.937956	251.907334
96	5.227328	274.649221
97	4.275410	217.232574
98	0.254191	65.519389
99	1.078914	86.139910

[100 rows x 2 columns]

```
[4]: #calling arrays
```

```
[5]: x= np.array(df["Age"]).reshape(-1,1)
```

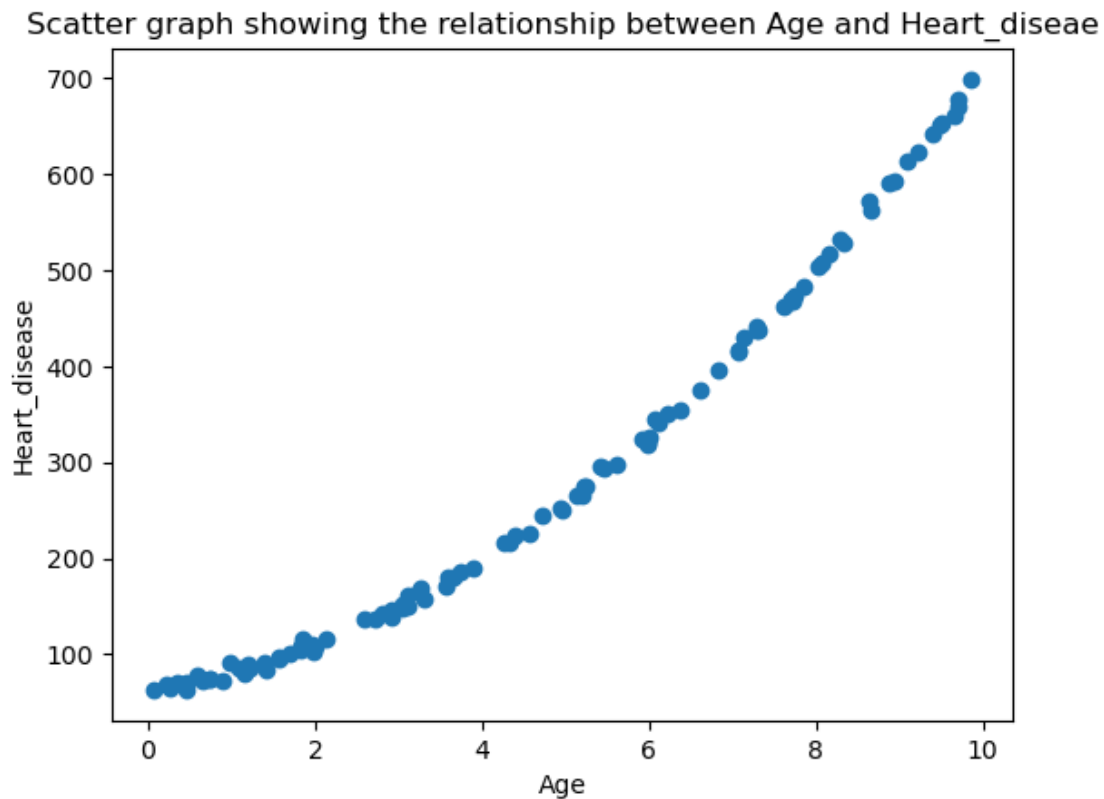
```
[6]: y=np.array(df["Heart_disease"])
```

```
[7]: #checking for missing data
df.isna().sum()
```

```
[7]: Age          0
Heart_disease    0
dtype: int64
```

```
[8]: #visualisation of the graph
```

```
[9]: import matplotlib.pyplot as plt
plt.scatter(x,y)
plt.xlabel("Age")
plt.ylabel("Heart_disease")
plt.title("Scatter graph showing the relationship between Age and Heart_disease")
#plt.grid(True)
plt.show()
```



```
[10]: #splitting data
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test =train_test_split(x,y,test_size = 0.2)
```

```
[11]: #standardizing data
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

```
[13]: #building model
from sklearn.linear_model import LinearRegression
model=LinearRegression()
```

```
model
```

```
[13]: LinearRegression()
```

```
[14]: #Model fitting  
model.fit(x_train_scaled,y_train)
```

```
[14]: LinearRegression()
```

```
[15]: #Making prediction  
y_pred=model.predict(x_test_scaled)  
y_pred
```

```
[15]: array([342.93280613, 195.83881853, 434.44797703, 533.01289795,  
          406.05101465, 367.76849487,  14.45988554, 467.30998378,  
          215.42854157, 187.34830324, 226.51718292, 186.89398322,  
          567.73619884, 331.34896842, 597.47981024, 374.46270551,  
          608.26641502, 117.50201731, 107.35942813, 199.30983018])
```

```
[16]: #getting coefficient  
model.coef_
```

```
[16]: array([186.74997827])
```

```
[17]: #Getting intercept  
model.intercept_
```

```
[17]: 276.658381441299
```

```
[18]: #Model accuracy on train values  
model.score(x_train_scaled,y_train)
```

```
[18]: 0.9623361610251643
```

```
[19]: #Model accuracy on test values  
model.score(x_test_scaled,y_test)
```

```
[19]: 0.9511348054806654
```

```
[20]: from sklearn.metrics import mean_absolute_error,r2_score,mean_squared_error  
mae=mean_absolute_error(y_test,y_pred)  
r2 = r2_score(y_test,y_pred)  
mse = mean_squared_error(y_test,y_pred)  
  
print(f"mae:{mae}")  
print(f"r2:{r2}")  
print(f"mse:{mse}")
```

```
mae:37.164980048735444
r2:0.9511348054806654
mse:1818.918837925957
```

## MODEL OPTIMIZATION

```
[21]: from sklearn.model_selection import GridSearchCV
      from sklearn.linear_model import LinearRegression
```

```
[37]: model=LinearRegression()
      model
```

```
[37]: LinearRegression()
```

```
[73]: #Define the parameter grid search
      param_grid = {
          'fit_intercept': [True, False],
          'copy_X': [True, False],
          'n_jobs': [True, False],

          }
      param_grid
```

```
[73]: {'fit_intercept': [True, False],
      'copy_X': [True, False],
      'n_jobs': [True, False]}
```

```
[74]: # Perform GridSearchCV
      grid_search = GridSearchCV(model, param_grid, cv=5)
      grid_search
```

```
[74]: GridSearchCV(cv=5, estimator=LinearRegression(),
                  param_grid={'copy_X': [True, False],
                              'fit_intercept': [True, False],
                              'n_jobs': [True, False]})
```

```
[75]: grid_search.fit(x_train_scaled, y_train)
      grid_search
```

```
[75]: GridSearchCV(cv=5, estimator=LinearRegression(),
                  param_grid={'copy_X': [True, False],
                              'fit_intercept': [True, False],
                              'n_jobs': [True, False]})
```

```
[78]: #Get the best parameters found from Grid search
      best_params = grid_search.best_params_
      best_params
```

```
[78]: {'copy_X': True, 'fit_intercept': True, 'n_jobs': True}
```

```
[82]: # Train LinearRegression model with the best parameters
best_model = LinearRegression(**best_params)
best_model.fit(x_train_scaled, y_train)
```

```
[82]: LinearRegression(n_jobs=True)
```

```
[85]: # Make predictions
y_pred_ridge = best_model.predict(x_test_scaled)
y_pred
```

```
[85]: array([342.93280613, 195.83881853, 434.44797703, 533.01289795,
        406.05101465, 367.76849487,  14.45988554, 467.30998378,
        215.42854157, 187.34830324, 226.51718292, 186.89398322,
        567.73619884, 331.34896842, 597.47981024, 374.46270551,
        608.26641502, 117.50201731, 107.35942813, 199.30983018])
```

```
[90]: # Model evaluation
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
print(f"mae:{mae}")
print(f"r2:{r2}")
print(f"mse:{mse}")
```

```
mae:37.164980048735444
r2:0.9511348054806654
mse:1818.918837925957
```